



# Before we start...

- This is technically a ‘members-only’ event (for student org. purposes.)
- Luckily, we count anyone in our Discord a member!

**<http://puhack.horse/discord>**



# Before we start...

Workshop structure:

- ~10 minutes: theory.
- 
-



# Before we start...

Workshop structure:

- ~10 minutes: theory.
- ~30 minutes: guided coding.
-



# Before we start...

Workshop structure:

- **~10 minutes:** theory.
- **~30 minutes:** guided coding.
- the rest of the time: self-directed hacking.



# Before we start...

Workshop structure:

- ~10 minutes: theory.
- ~30 minutes: guided coding.
- the rest of the time: self-directed hacking.

**guided coding: you'll write a matrix multiplication kernel.**



# Before we start...

Workshop structure:

- ~10 minutes: theory.
- ~30 minutes: guided coding.
- the rest of the time: self-directed hacking.

**guided coding:** you'll write a matrix multiplication kernel.

**hacking:** write some other kernels. maybe convolution.



# Before we start...

Hi! Rules for this presentation:

- You *should* ask questions during the presentation.
-



# Before we start...

Hi! Rules for this presentation:

- You *should* ask questions during the presentation.
- in fact...

**Ask me a question right now.**



# Before we start...

Hi! Rules for this presentation:

- You *should* ask questions during the presentation.
- in fact...

**Ask me a question right now.**

**Another.**



# Introduction to GPU Programming

Kartavya Vashishtha

February 26, 2026

# A GPU from a thousand feet away



# A GPU from a thousand feet away

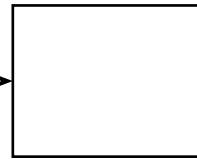


where is it?!

# A GPU from a thousand feet away



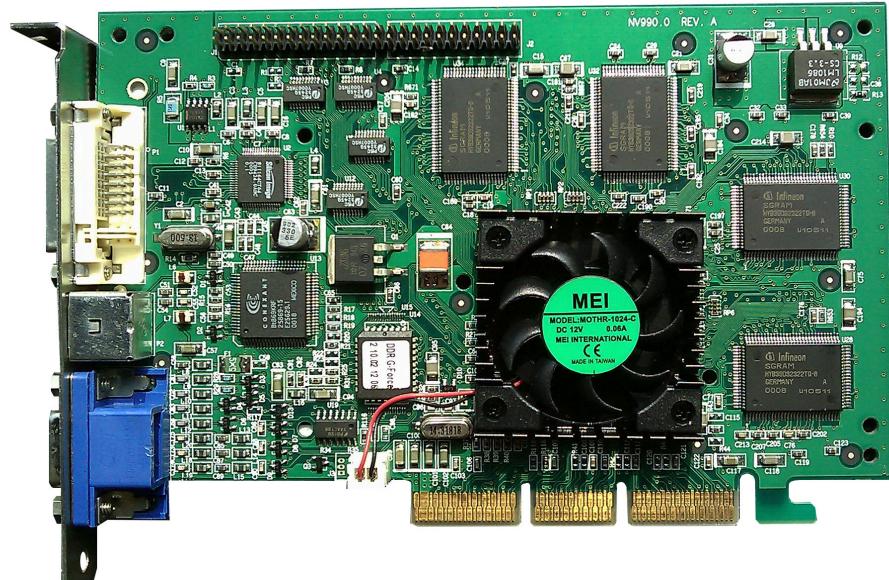
oh need to zoom in →



# A GPU from a few inches away



- not a CPU
- 



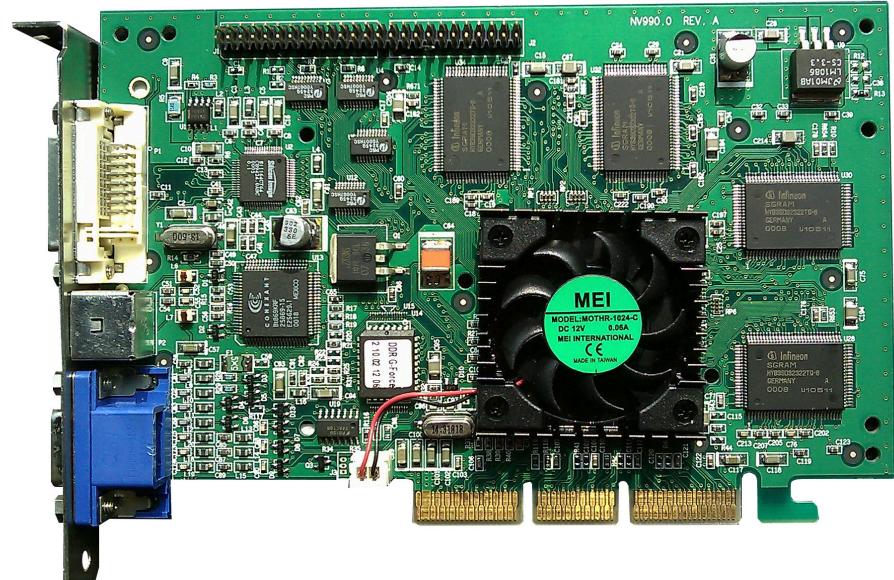
120 MHz Nvidia GeForce 256

# A GPU from a few inches away



- not a CPU
- trades single-thread speed for many-thread throughput

(more on this later!)



120 MHz Nvidia GeForce 256



# A History of GPU Usage

1. **Graphics era:** Fixed-function rendering. Push pixels to a buffer, modify, and dump to screen. Rigid APIs.



# A History of GPU Usage

1. **Graphics era:** Fixed-function rendering. Push pixels to a buffer, modify, and dump to screen. Rigid APIs.
2. **Mid-2000s:** *what if... we use these chips for things that aren't games.* NVIDIA releases CUDA (2007) — general-purpose computing API.

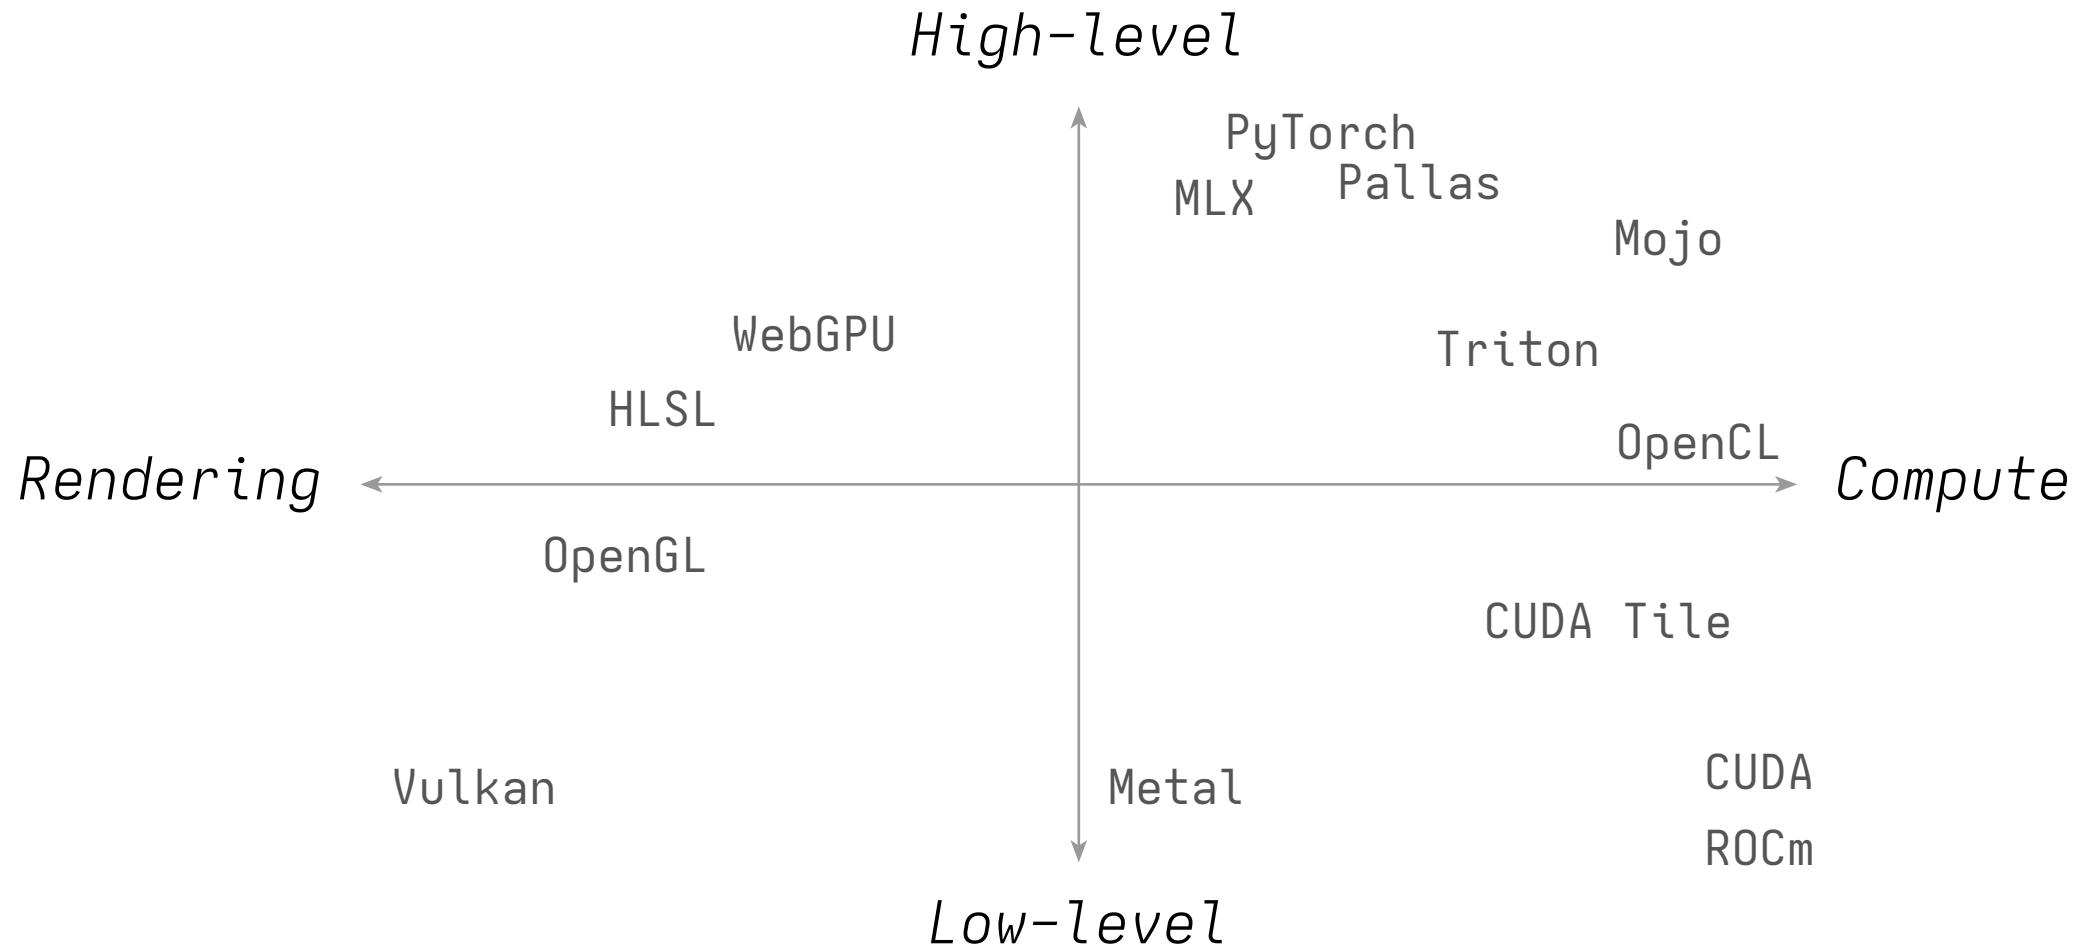


# A History of GPU Usage

1. **Graphics era:** Fixed-function rendering. Push pixels to a buffer, modify, and dump to screen. Rigid APIs.
2. **Mid-2000s:** *what if... we use these chips for things that aren't games.* NVIDIA releases CUDA (2007) — general-purpose computing API.
3. **2010s:** people start running machine learning on them.

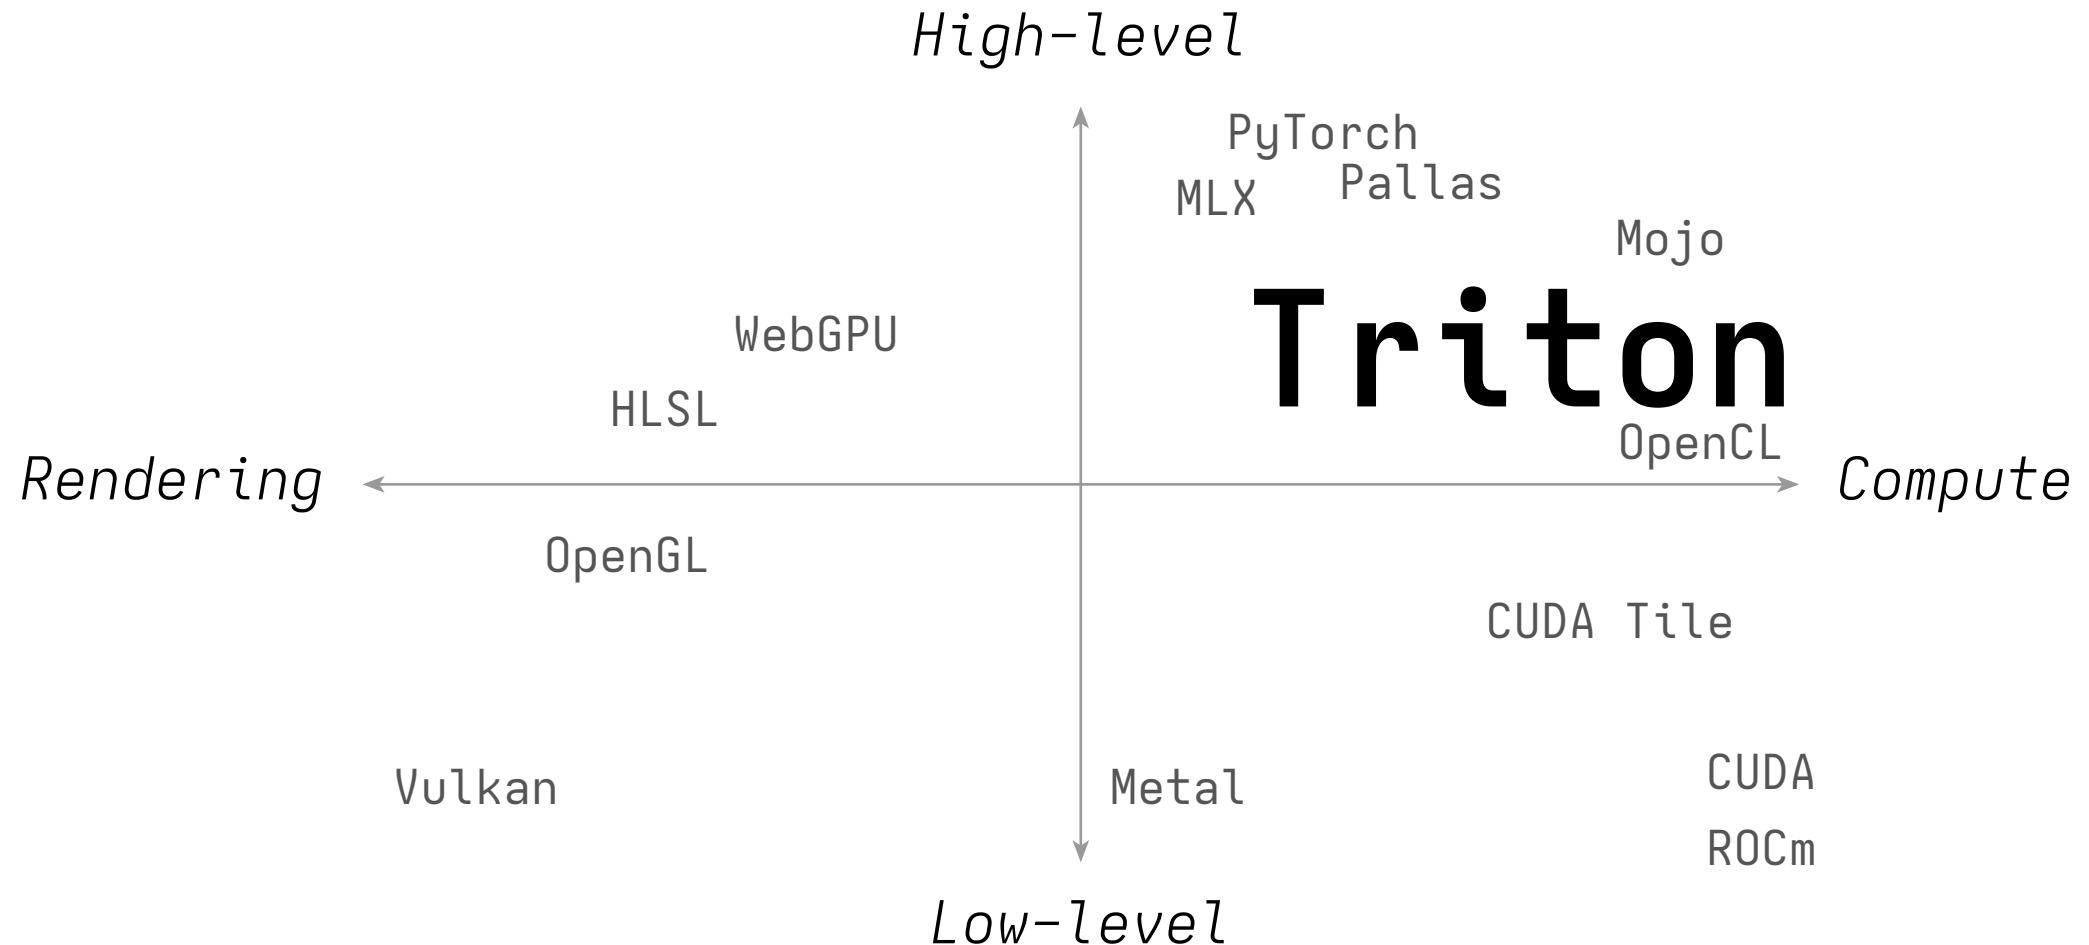


# Many ways to program a GPU



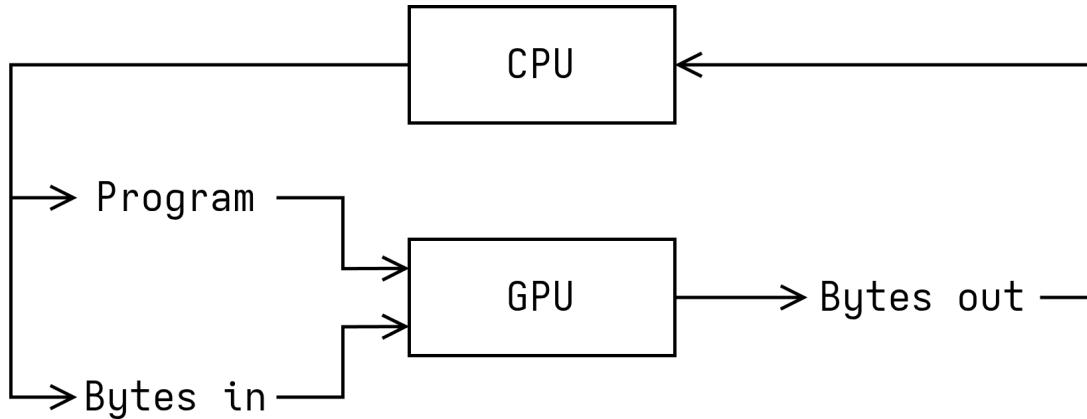


# Many ways to program a GPU



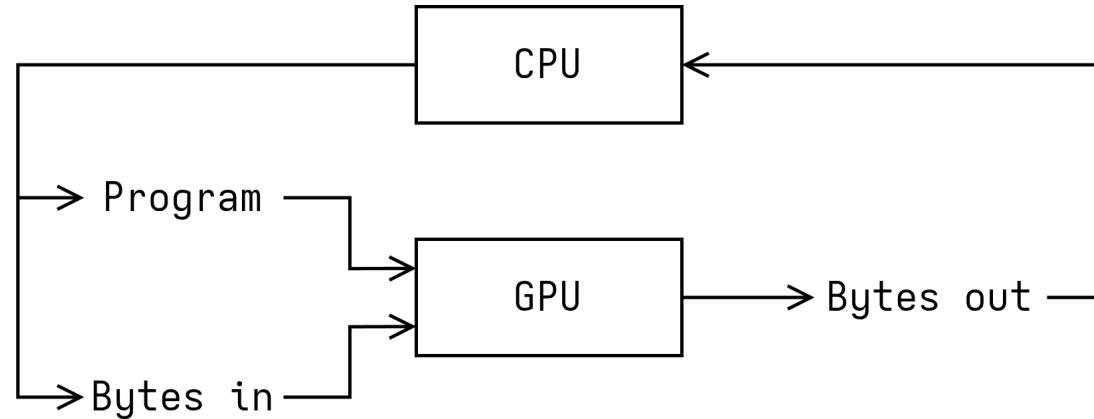


# A simple model





# An (overly) simple model





# Switching gears: time to write ReLU!

*Recall, or learn very quickly,*

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

Pseudocode:



# Switching gears: time to write ReLU!

*Recall, or learn very quickly,*

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

Pseudocode:

```
def relu(x):  
    return max(x, 0)
```



# Switching gears: time to write ReLU!

*Recall, or learn very quickly,*

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

Pseudocode:

```
def relu(x):  
    return max(x, 0)
```

For an array:



# Switching gears: time to write ReLU!

*Recall, or learn very quickly,*

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{else} \end{cases}$$

Pseudocode:

```
def relu(x):
    return max(x, 0)
```

For an array:

```
def relu_array(arr):
    return [max(x, 0) for x in arr]
```



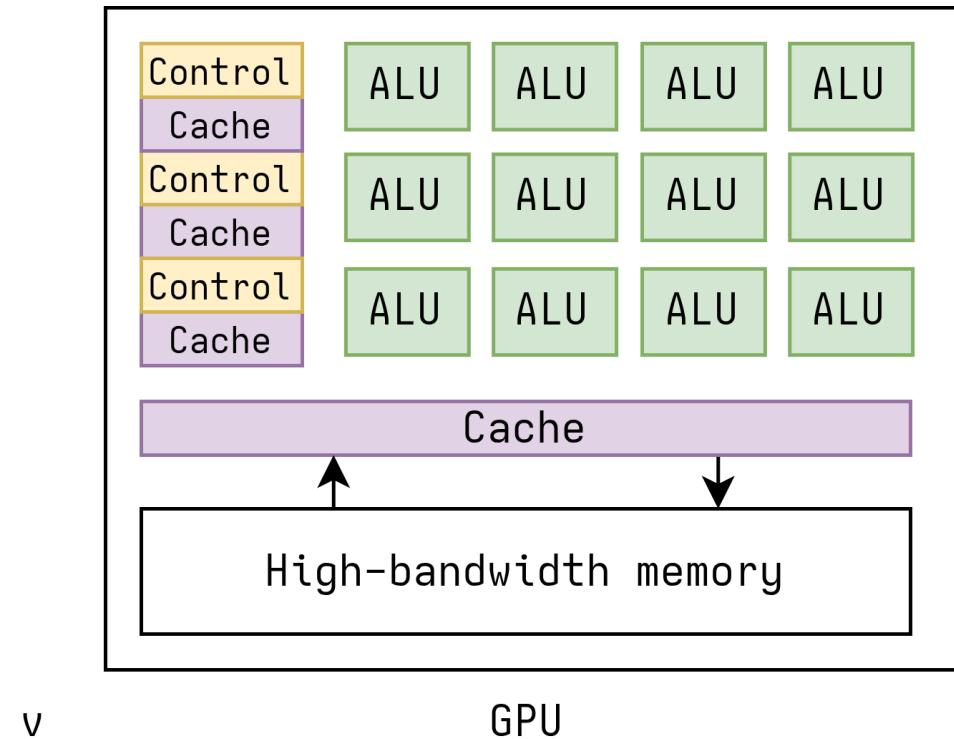
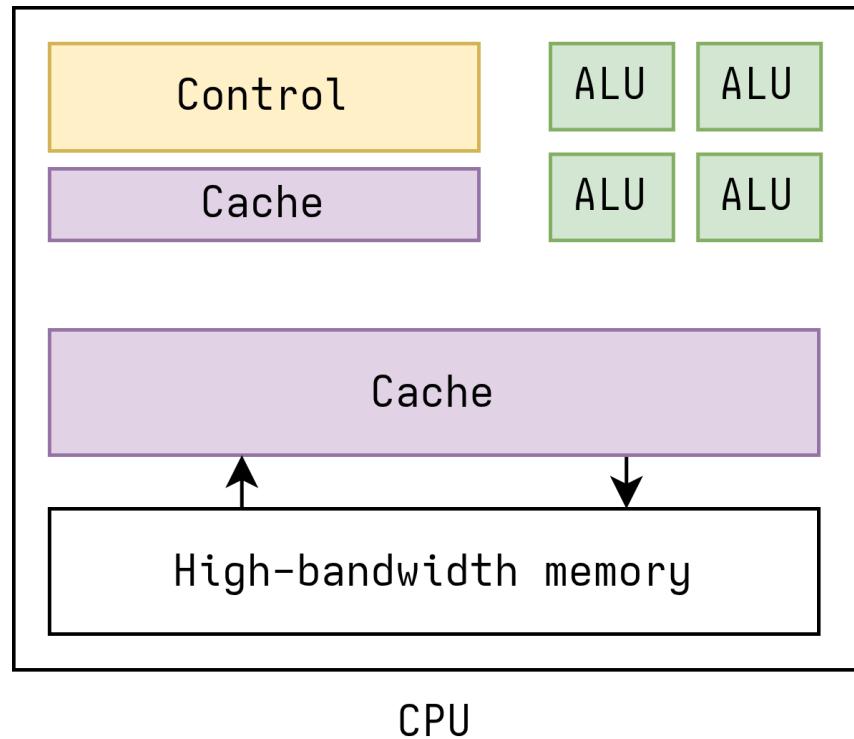
# A GPU from a few millimeters away

- trades single-thread speed for many-thread throughput
- question: *how do you divide work between these threads?*



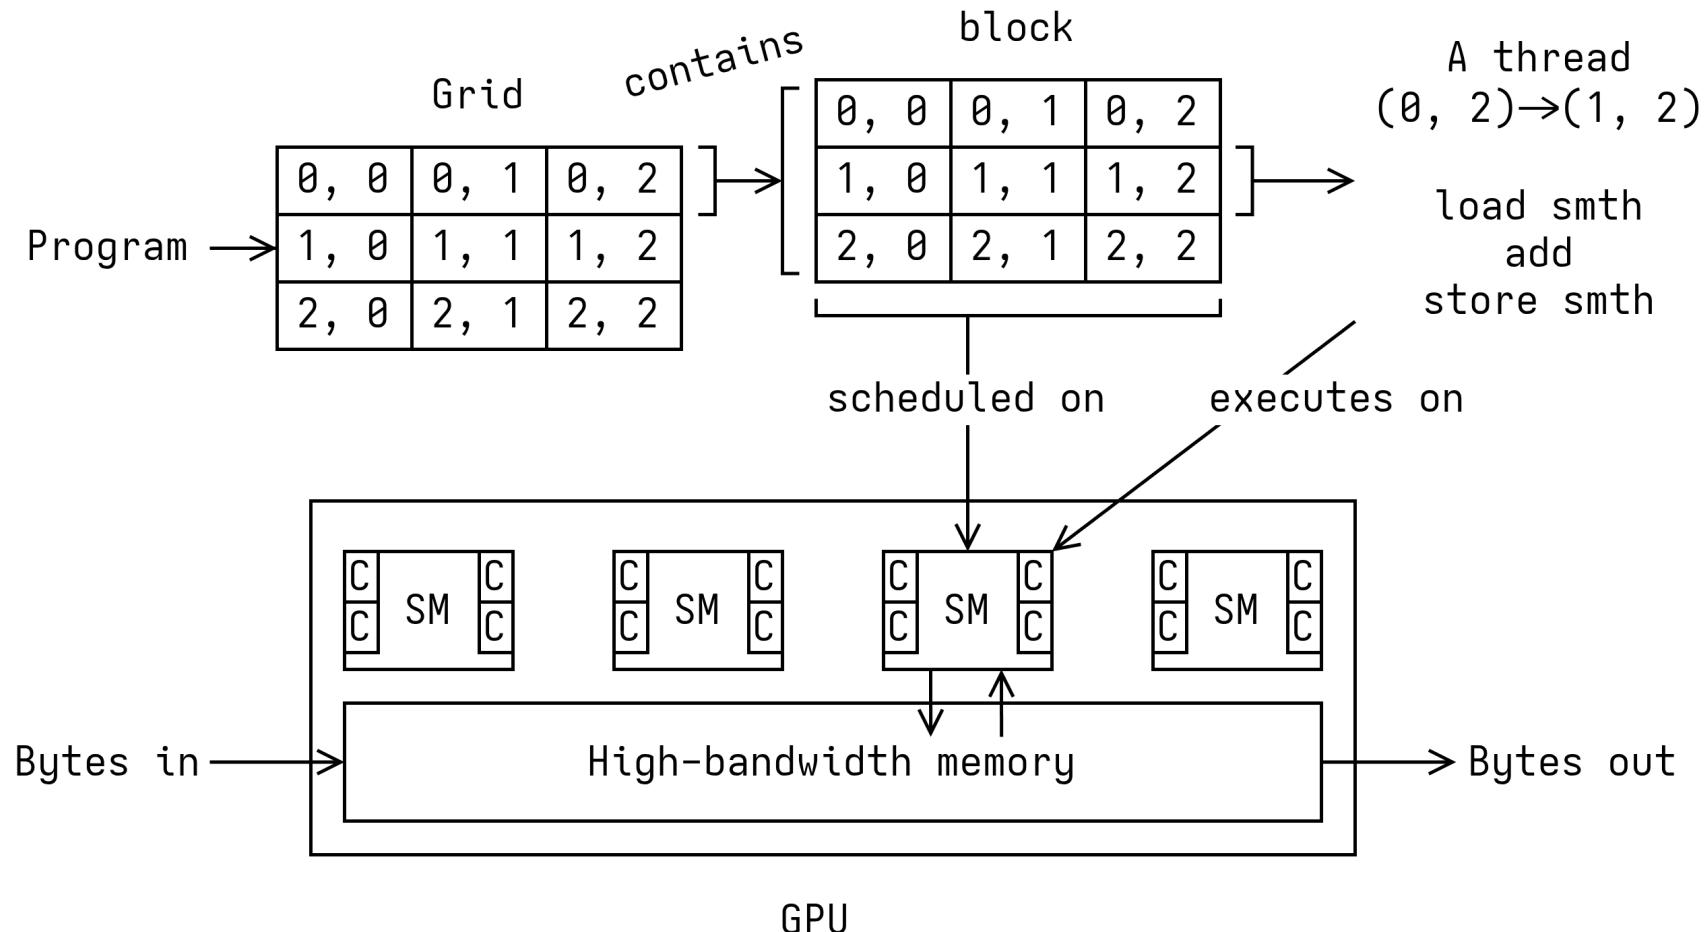
# A GPU from a few millimeters away

- trades single-thread speed for many-thread throughput
- question: *how do you divide work between these threads?*





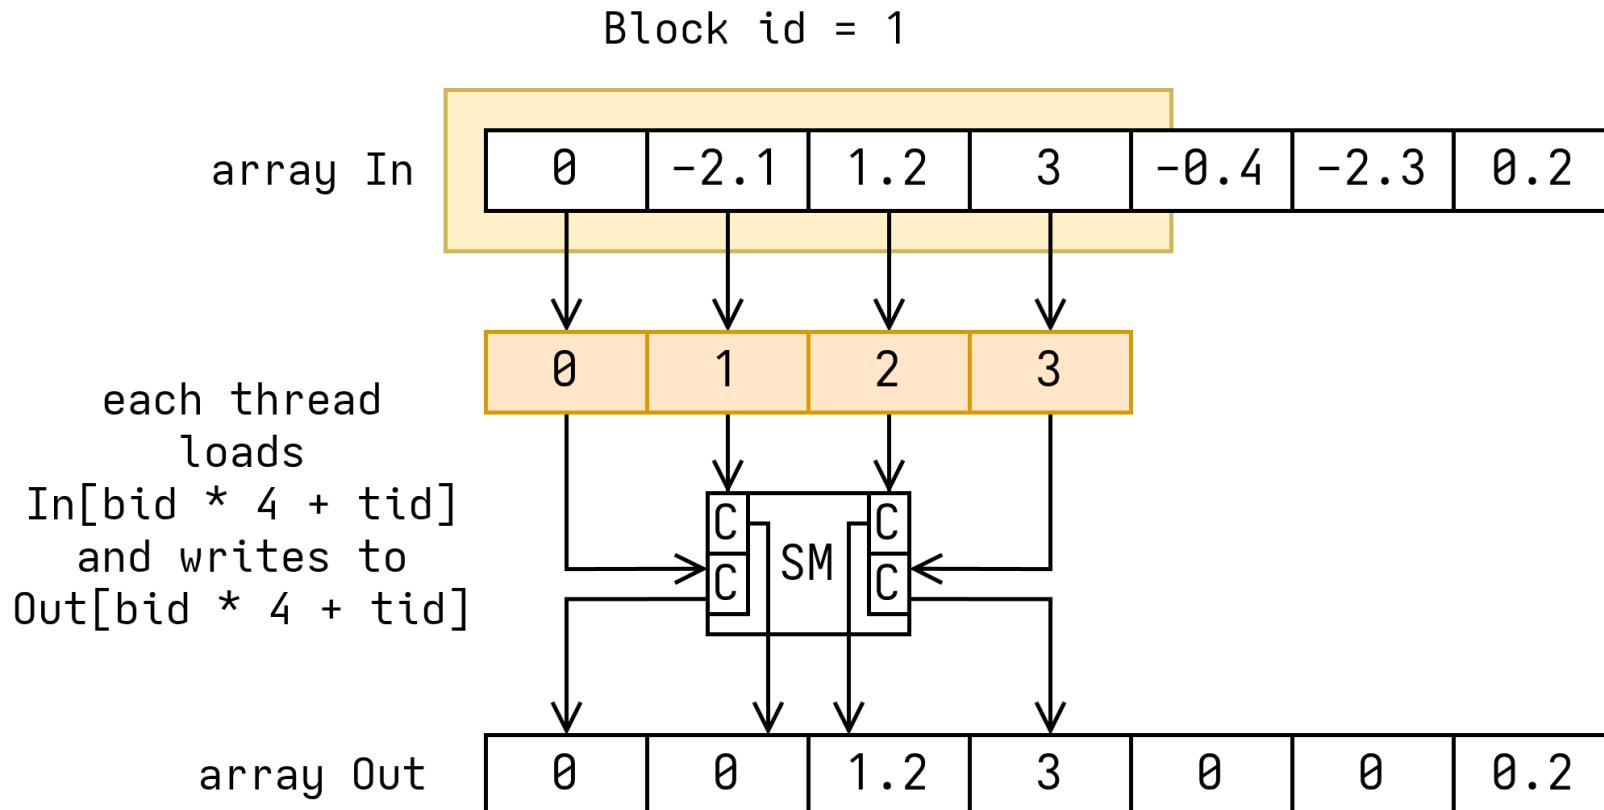
# A simple model for a GPU





# Taking ReLU to the GPU

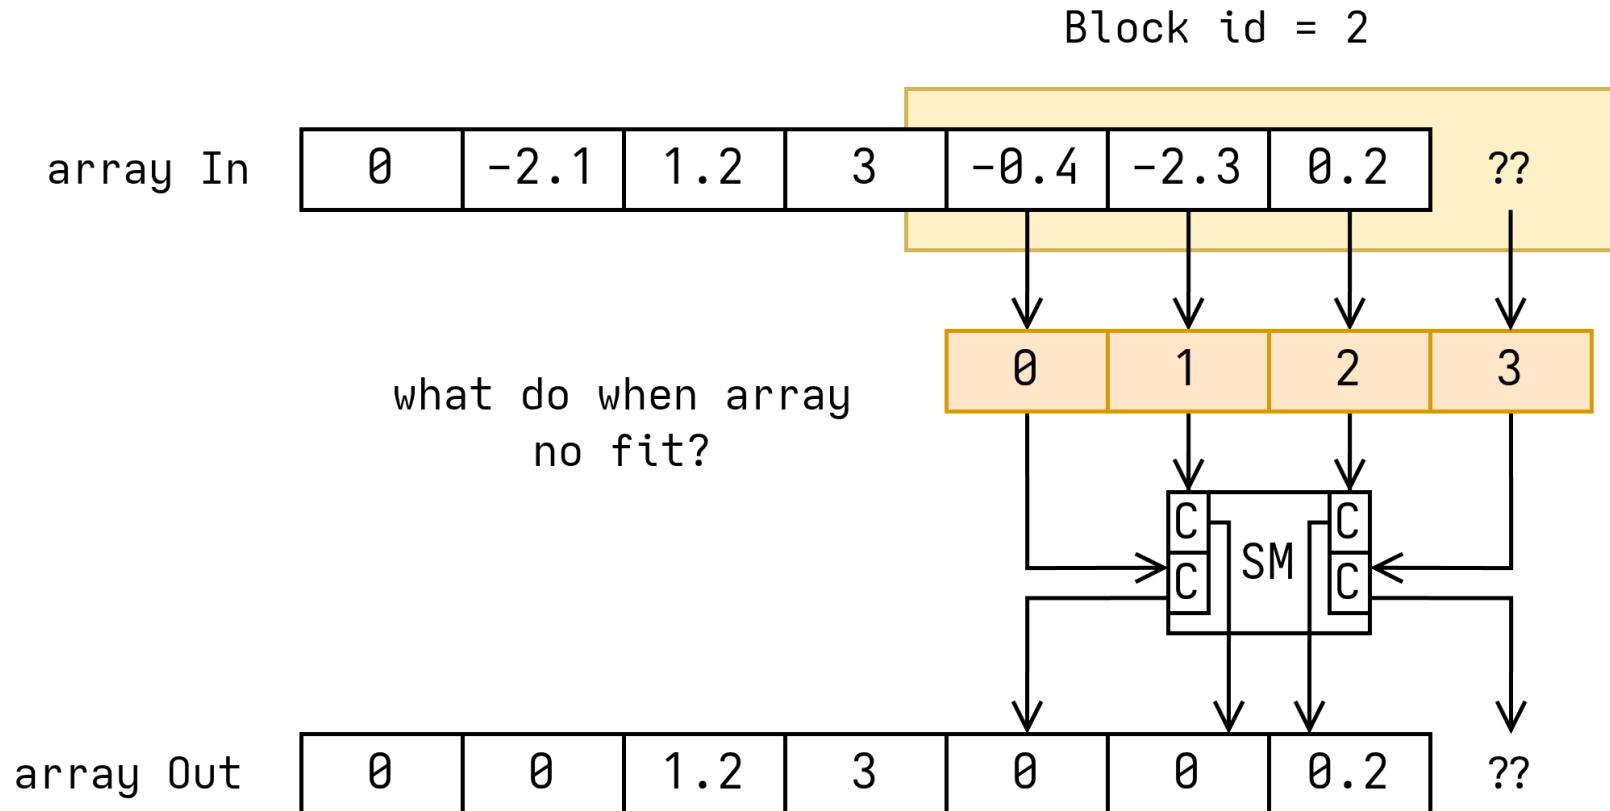
Grid with dimension (2); blocks are [0, 1]. Each block has 4 threads.





# Taking ReLU to the GPU

What happens at block 1?





# Time to move over to writing code!!

- go to <https://puhack.horse/gpu-colab>
- click ‘Copy to Drive’ in top bar.



# An example problem decomposition

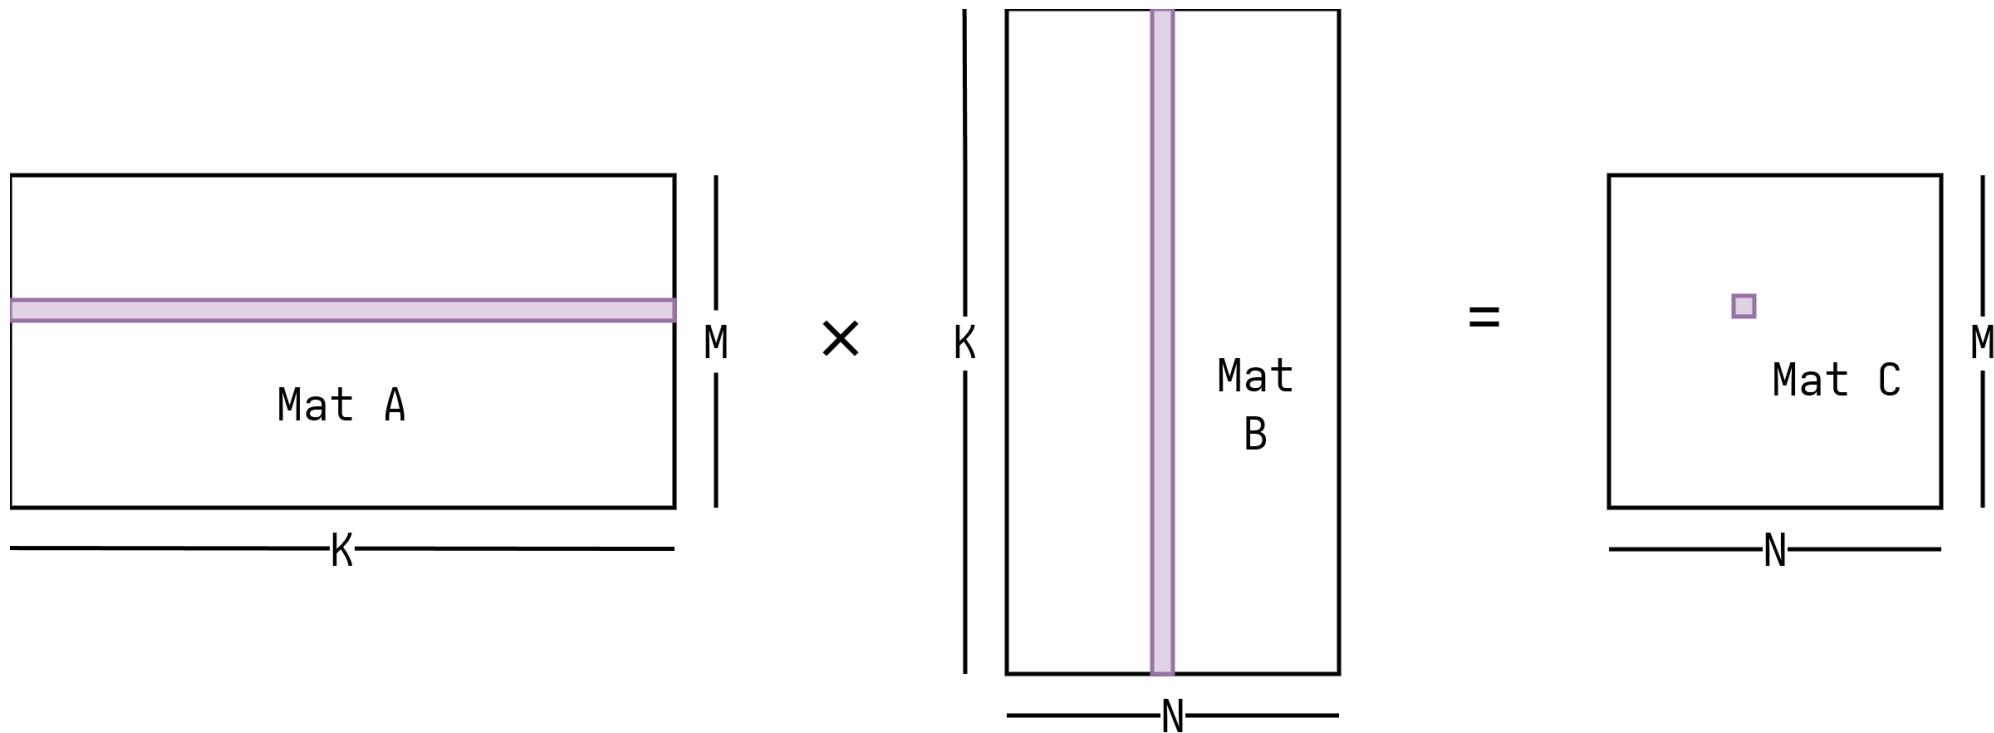
Hm. What sort of problems can be broken into parallel chunks?



# An example problem decomposition

Hm. What sort of problems can be broken into parallel chunks?

Enter the humble matmul.





# The humble matmul

```
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

- Imagine one thread per  $(i, j)$  index:
  - ▶ Thread  $(0, 0)$  loads  $A[0, k]$  and  $B[k, 0]$
  - ▶ Thread  $(0, 1)$  loads  $A[0, k]$  and  $B[k, 1]$
  - ▶ Thread  $(1, 0)$  loads  $A[1, k]$  and  $B[k, 0]$
  - ▶ etc.
- 
- 
- 
-



# The humble matmul

```
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

- Imagine one thread per (i, j) index:
  - ▶ Thread (0, 0) loads A[0, k] and B[k, 0]
  - ▶ Thread (0, 1) loads A[0, k] and B[k, 1]
  - ▶ Thread (1, 0) loads A[1, k] and B[k, 0]
  - ▶ etc.
- GPUs execute a whole block on one compute unit.
- 
- 
-



# The humble matmul

```
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

- Imagine one thread per  $(i, j)$  index:
  - ▶ Thread  $(0, 0)$  loads  $A[0, k]$  and  $B[k, 0]$
  - ▶ Thread  $(0, 1)$  loads  $A[0, k]$  and  $B[k, 1]$
  - ▶ Thread  $(1, 0)$  loads  $A[1, k]$  and  $B[k, 0]$
  - ▶ etc.
- GPUs execute a whole block on one compute unit.
- This compute unit *has very fast memory*.
- 
-



# The humble matmul

```
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

- Imagine one thread per  $(i, j)$  index:
  - ▶ Thread  $(0, 0)$  loads  $A[0, k]$  and  $B[k, 0]$
  - ▶ Thread  $(0, 1)$  loads  $A[0, k]$  and  $B[k, 1]$
  - ▶ Thread  $(1, 0)$  loads  $A[1, k]$  and  $B[k, 0]$
  - ▶ etc.
- GPUs execute a whole block on one compute unit.
- This compute unit *has very fast memory*.
- Sharing these accesses is *good*.
-

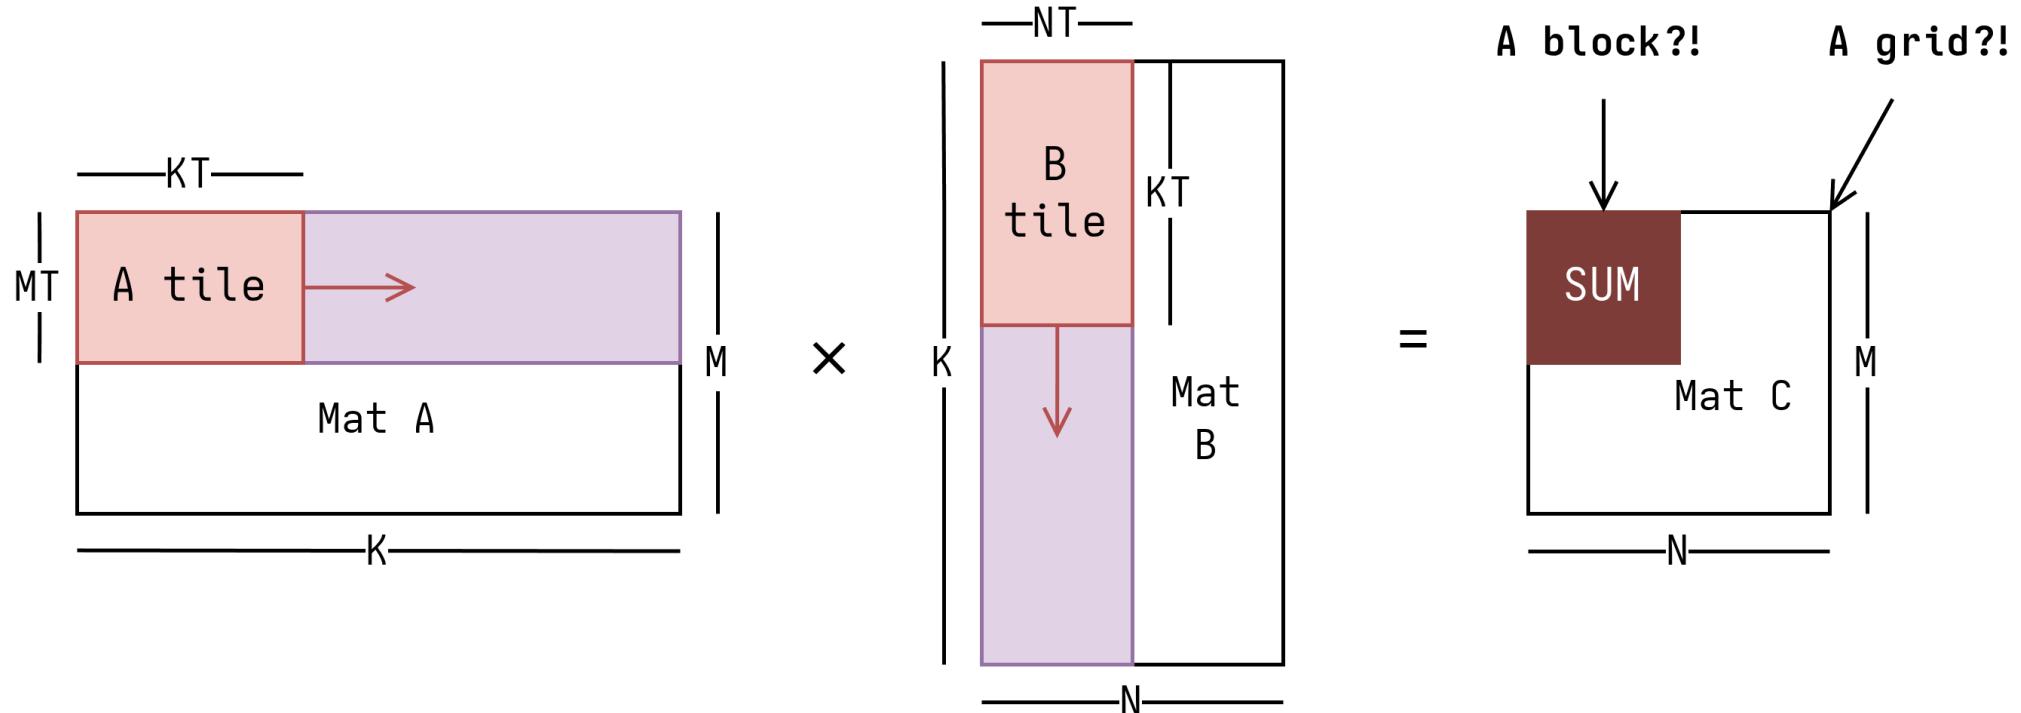
# The humble matmul



```
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

- Imagine one thread per  $(i, j)$  index:
  - ▶ Thread  $(0, 0)$  loads  $A[0, k]$  and  $B[k, 0]$
  - ▶ Thread  $(0, 1)$  loads  $A[0, k]$  and  $B[k, 1]$
  - ▶ Thread  $(1, 0)$  loads  $A[1, k]$  and  $B[k, 0]$
  - ▶ etc.
- GPUs execute a whole block on one compute unit.
- This compute unit *has very fast memory*.
- Sharing these accesses is *good*.
- So let's write this with tiles!

# Presenting: the tiled matmul





# Tiled matmul pseudocode

```
def matmul_kernel(A, B, C, M, N, K):
    c_m, c_n = program_id(0), program_id(1)
    accumulator = zeros((MT, NT))

    for k in range(0, K, KT):
        tile_A = A[c_m * MT : (c_m + 1) * MT,           k : k + KT      ]
        tile_B = B[          k : k + KT,           c_n * NT : (c_n + 1) * NT]

        accumulator += dot(tile_A, tile_B)

    C[c_m * MT : (c_m + 1) * MT, c_n * NT : (c_n + 1) * NT] = accumulator
```



# Tiled matmul pseudocode

```
def matmul_kernel(A, B, C, M, N, K):
    c_m, c_n = program_id(0), program_id(1)
    accumulator = zeros((MT, NT))

    for k in range(0, K, KT):
        tile_A = A[c_m * MT : (c_m + 1) * MT,           k : k + KT]
        tile_B = B[          k : k + KT,           c_n * NT : (c_n + 1) * NT]

        accumulator += dot(tile_A, tile_B)

    C[c_m * MT : (c_m + 1) * MT, c_n * NT : (c_n + 1) * NT] = accumulator
```

**CODING TIEM!! ← so excited I cna't spell**



# Hacking Section!

You can:

- go to <https://tensara.org/>
- or
- go to <https://puhack.horse/gpu-colab-conv>